

Score your own AI agent

1 Run the conversation with your agent.

2 Tick every check the transcript passes. No proof in the transcript, no tick.

3 Add the points you ticked. That is your score out of 100 for that lane.

Shopping Assistant

Pre-sale, ends in a real recommendation

Support Agent

Post-sale, resolve without a human

Did it answer you30

☐ Answered the question. Direct, substantive reply to each question. A generic greeting or boilerplate fails.14

☐ Didn't contradict itself. No contradictions across turns, no invented policies or specs.9

☐ Left nothing unanswered. No question left unanswered or answered with something unrelated.7

Did it work out what you needed20

☐ Asked before recommending. Asks a relevant clarifying question on a vague opener before recommending.8

☐ Built on your answers. Uses the shopper's answers to refine across turns, builds a profile.7

☐ Didn't dump a product list. Does not dump a generic best-seller list before understanding needs.5

Did it recommend well22

☐ Named a specific product. Recommends at least one specific, named product.9

☐ Explained why it fits. Ties the pick to the shopper's stated constraints with real reasoning.8

☐ Suited the store and the ask. Recommendations fit the store's catalog and the request.5

Did it show you the product18

Code-checked, not judged. These four pass only if the link, price, review, or comparison actually appears in the transcript.

Did it solve the problem40

All three fail if it sent you away. Pushing you to "email us" or "call us" on an answerable question scores this whole group zero.

☐ Answered it in the chat. Primary ask gets a complete, store-specific answer in-channel. A generic pointer fails.18

☐ Ended the right way. Outcome-correct ending: resolved, or a justified handover done well. An unjustified bail fails.12

☐ Didn't send you away. Does not push out of channel ("email us") when the question was answerable in chat.10

Was it accurate25

☐ Gave real facts. At least two concrete, store-specific facts. One vague reassurance fails.13

☐ Sounded like this store. Claims read as grounded in this store, not industry generalities.7

☐ Didn't contradict itself. No self-contradiction across turns.5

Could you act on it20

☐ Told you what to do. Customer leaves with steps they can execute now, in a cold session. "Reach out if..." alone fails.12

☐ Answered all of it. Covers the actual ask, not a fragment.8

Did it close properly15

☐ Said what happens next. Sets expectations on who acts and when. Vague "soon" fails.8

☐ Finished cleanly. Clean, complete close.7

Your Support Agent score

/ 100

☐	Gave a link to open. An openable product link or card.	7
☐	Showed a price. A concrete price attached to a recommended product.	6
☐	Backed it with reviews. Review counts or ratings support the recommendation.	3
☐	Offered a comparison. Multiple distinct options laid out to compare.	2
Did it help you buy		10
☐	Gave a next step. Ends with a concrete next step.	5
☐	Completed the purchase. Completes add-to-cart or checkout when explicitly asked.	3
☐	Finished cleanly. Closes cleanly, no dangling question or mid-thought stop.	2

Your Shopping Assistant score
/ 100

Fairness rules baked in

- **A good handover is not a failure.** Containment is measured separately by automation rate, so quality never double-penalizes it.
- **Judge only what a first-time visitor could see.** Cold session, no account. Never penalized for not knowing your order history.
- **Broken widgets don't count as zero.** A conversation needs at least three cleanly-timed answers to count. Dropped connections are store transport, so they're excluded, not scored.
- **Every run is audited.** An adversarial second pass re-checks sampled verdicts. A run is trusted only at 90% agreement or higher.

This rubric is the **quality** score. A vendor's lane ranking also weighs automation and speed. **Shopping** 40 / 35 / 25 · **Support** 50 / 40 / 10. evals.gorgias.com