

The rubric

Every conversation is scored out of 100 by a blind LLM judge against fixed, binary checks. Each passing check must quote the transcript. The judge never picks a number; the score is derived from the checks. Two rubrics, one per job.

PRE-SALE WEIGHTING		POST-SALE WEIGHTING	
Answered you	30	Solved the problem	40
Recommended well	22	Was accurate	25
Worked out your need	20	Could act on it	20
Showed the product	18	Closed properly	15
Helped you buy	10		

Pre-sale · 16 checks in 5 groups

CHECK	PASSES WHEN	POINTS
Did it answer you		30
Answered the question	Every question got a real reply. A greeting or brand waffle does not count.	14
Didn't contradict itself	No conflicting claims, and no policies or specs it made up.	9
Left nothing unanswered	No question skipped or answered with something unrelated.	7
Did it work out what you needed		20
Asked before recommending	On a vague opener, found out what the shopper needed first.	8
Built on your answers	Used what the shopper said to narrow things down, instead of repeating a generic pick.	7
Didn't dump a product list	No wall of best-sellers before understanding the need.	5
Did it recommend well		22
Named a specific product	Recommended something real, not a category.	9
Explained why it fits	Tied the recommendation to what the shopper actually said they needed.	8
Suited the store and the ask	The pick makes sense for that catalog and that request.	5
Did it show you the product		18
Gave a link to open	A product link or card the shopper can click.	7

CHECK	PASSES WHEN	POINTS
Showed a price	A real price attached to the recommendation.	6
Backed it with reviews	Ratings or review counts supporting the pick.	3
Offered a comparison	More than one option laid out side by side.	2
Did it help you buy		10
Gave a next step	Ended with something concrete to do.	5
Completed the purchase	Actually added to cart or checked out when asked to.	3
Finished cleanly	No dangling question, no mid-thought stop.	2

Post-sale · 10 checks in 4 groups

CHECK	PASSES WHEN	POINTS
Did it solve the problem		40
Answered it in the chat	A complete answer specific to that store. Pointing at a help page does not count.	18
Ended the right way	Sorted it, or handed to a person properly with the details collected first.	12
Didn't send you away	No "email us" or "call us" when it could have just answered.	10
Was it accurate		25
Gave real facts	At least two concrete details: a number, a timeframe, a named policy.	13
Sounded like this store	The claims read as that shop's actual policy, not generic advice.	7
Didn't contradict itself	No conflicting claims across the conversation.	5
Could you act on it		20
Told you what to do	Steps you can carry out right now, without an account.	12
Answered all of it	Covered the whole question, not just part.	8
Did it close properly		15
Said what happens next	Named who acts and when. "Soon" does not count.	8
Finished cleanly	Nothing left hanging.	7

Worked example · Conversation #1841 · Pre-sale lane

How a single conversation scores, check by check.

VERDICT	CHECK	EVIDENCE	POINTS
Pass	Product fact match	"Order #88421 is in our warehouse" matches order-system data.	5
Pass	Policy accuracy	"5-7 days standard" matches the published policy page.	5
Pass	No hallucinated SKUs	No product identifiers or variant names invented.	5
Pass	Price precision	No price stated, so no penalty.	5
Fail	Query fully addressed	"What can I do?" got no actionable path for the deadline.	0
Pass	Actionable next step	Contact again after shipment. Minimal but acceptable.	3
Pass	No over-apology	Apology appears once, attached to a real error.	3
Pass	Scannable format	Short paragraphs, information chunked clearly.	4
Pass	Sentence economy	All sentences under 25 words, no padding.	3
Pass	Escalation threshold	Did not escalate a routine shipping-status query.	3
Accuracy 20 · Completeness 8 · Tone 7 · Structure 7			36

Gorgias Ecommerce AI Agent Benchmark · Rubric v2.3 · Full pipeline, judge prompts, and inter-rater reliability report available on request. Every score in the report traces back to these checks.